

Shadoks Approach to Parallel Reconfiguration of Triangulations

Guilherme D. da Fonseca  

LIS, Aix-Marseille Université

Fabien Feschet  

LIMOS, Université Clermont Auvergne

Yan Gerard  

LIMOS, Université Clermont Auvergne

1 Abstract

We describe the methods used by Team Shadoks to win the CG:SHOP 2026 Challenge on parallel reconfiguration of planar triangulations. An instance is a collection of triangulations of a common point set. We must select a center triangulation and find short parallel-flip paths from each input triangulation to the center, minimizing the sum of path lengths. Our approach combines exact methods based on SAT with several greedy heuristics, and also makes use of SAT and MaxSAT for solution improvement. We present a SAT encoding for bounded-length paths and a global formulation for fixed path-length vectors. We discuss how these components interact in practice and summarize the performance of our solvers on the benchmark instances.

2012 ACM Subject Classification Theory of computation → Computational geometry

Keywords and phrases Heuristic, exact algorithm, SAT solver, MaxSAT solver, computational geometry

Category CG Challenge

Related Version arxiv.org/abs/

Supplementary Material (*Source Code*): <https://github.com/gfonsecabr/shadoks-CGSHOP2026>

Acknowledgements We would like to thank the Challenge organizers and other competitors for their time, feedback, and making this whole event possible. We would like to thank Hélène Toussaint, Raphaël Amato, Boris Lonjon, and William Guyot-Lénat from LIMOS, as well as the Qarma and TALEP teams and Manuel Bertrand from LIS, who continue to make the computational resources of the LIMOS and LIS clusters available to our research. We would also like to thank Aldo Gonzalez-Lorenzo for the very useful discussion on SAT models.

10 1 Introduction

The CG:SHOP Challenge is an annual competition in geometric optimization. In its eighth edition in 2026, the challenge focuses on a reconfiguration problem between planar triangulations. Our team, called *Shadoks*, won first place with the best solution (among the 28 participating teams) to 249 instances out of 250 instances and provably optimal solutions to 189 instances.

In this paper, we outline the exact methods and heuristics that we employed. We start with some definitions that allow us to describe the problem. Throughout, we consider triangulations of a common point set $S \subset \mathbb{R}^2$. Given a triangulation T , a *unit flip* is the operation that removes an edge $e \in T$ and adds an edge e' , obtaining a new triangulation $T' = T \setminus \{e\} \cup \{e'\}$. Notice that the edge e must cross e' . Similarly, a *parallel flip* removes a set of edges $E \subset T$ and adds a set of edges E' , in a way that $T' = T \setminus E \cup E'$ is a triangulation, with the condition that no two edges of E are in the same triangle in T . A *path* of length ℓ is

25 a sequence of triangulations T_0, \dots, T_ℓ such that for all i , the triangulation T_{i+1} is obtained
 26 from T_i by performing a parallel flip.

27 An *instance* is a set $S \subset \mathbb{R}^2$ of n points and a set $\mathcal{T} = \mathbf{T}_1, \dots, \mathbf{T}_{|\mathcal{T}|}$ of triangulations of
 28 S , called *input triangulations*. A *solution* is a set of paths $P_1, \dots, P_{|\mathcal{T}|}$ such that P_i starts at
 29 \mathbf{T}_i for all i and all paths end in a common triangulation called *center*. The goal is to find a
 30 solution that minimizes the *objective value* defined as the sum of the lengths of its paths.

31 During the competition, the organizers provided a total of 250 instances, with n ranging
 32 from 15 to 12,500 points and $|\mathcal{T}|$ ranging from 2 to 200 triangulations. The 250 instances are
 33 divided into three classes: 100 `random` instances, 101 `woc` instances, and 49 `rirs` instances.
 34 The former two instances have up to 320 points and 2 to 20 input triangulations (hence,
 35 we call them `small` instances), while the latter have 500 to 12500 points and 20 to 200
 36 input triangulations. The centers of some of our best solutions are presented in Figure 1.
 37 Additional details about the challenge can be found in the organizers' survey paper [5].

38 Our best solvers heavily rely on the SAT solver `CaDiCa1` [3] and the MaxSAT solver
 39 `EvalMaxSAT` [2]. Nevertheless, we also developed heuristics that do not rely on any external
 40 solver, which are important to find initial solutions to some large instances, which are then
 41 improved by roughly 10% using SAT and MaxSAT solvers. Furthermore, we managed to
 42 solve 189 of the 201 `small` instances exactly by repeatedly using the SAT solver as well as
 43 some lower bounds.

44 Mention the other teams strategy here...

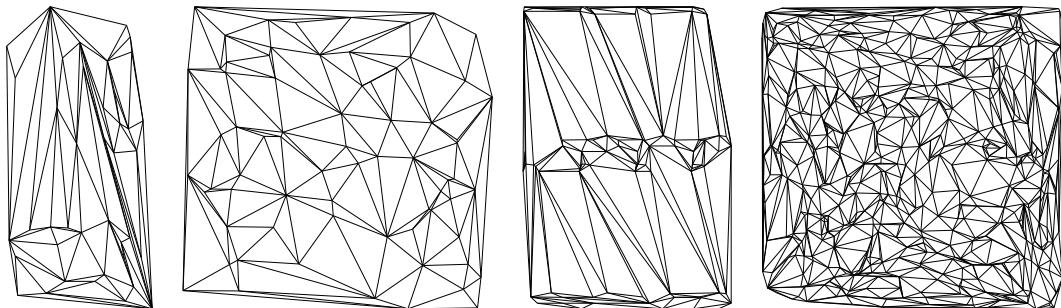
45 We describe our exact algorithms in Section 2, the heuristics in Section 3, and discuss
 46 the results we obtained in Section 4. Concluding remarks and open problems are presented
 47 in Section 5.

48 2 Exact Algorithms

49 This section describes all elements of our exact solver, many of which are also used in the
 50 heuristic solvers. We first show how to use a SAT solver to compute shortest paths between
 51 two triangulations (Section 2.1). We then show how to extend this result to test if a solution
 52 with a list of path lengths exists (Section 2.2). We show how to obtain lower bounds in
 53 Section 2.3 and put the previous elements together to describe our exact solver in Section 2.4.

54 2.1 Path SAT Formulation


55 We now describe a SAT formulation for the following decision problem. The input is a set S
 56 of n points, an integer ℓ and two triangulations T_0, T_ℓ . The output is whether there exists a



11  **Figure 1** Our best centers to instances `random_78_40_10`, `woc-70-random-9a7d18d3`,
 12 `woc-90-tsplib`, and `rirs-500-50-23d00ec5`, respectively.

57 path T_0, \dots, T_ℓ of length ℓ .

58 We define two types of variables. For $i = 0, \dots, \ell$ and for $u \neq v \in S$, we define an *edge*
 59 *variable* $e(u, v, i)$. The variable $e(u, v, i)$ represents that the edge uv is in the triangulation T_i .
 60 There are $O(n^2\ell)$ edge variables. It would be possible to define a SAT formulation using only
 61 such variables. However, a SAT formulation that performed much better in our experiments
 62 uses a second type of variable.



63 **Figure 2** Illustration of a flip variable $f(u, v, u_2, v_2, i)$.

64 We say that a convex quadrilateral is *empty* if it contains no point of S except for its
 65 vertices. For $i = 0, \dots, \ell$ and for $u \neq v \neq u_2 \neq v_2 \in S$ such that u, u_2, v, v_2 form an empty
 66 convex quadrilateral, we define a *flip variable* $f(u, v, u_2, v_2, i)$. The variable $f(u, v, u_2, v_2, i)$
 67 represents a unit flip such that the edge uv is in triangulation T_i and u_2v_2 is in triangulation
 68 T_{i+1} , as shown in Figure 2. Notice that if the points are uniformly distributed, then
 69 the number of empty convex quadrilaterals is $\Theta(n^2)$ [4], which means that for uniformly
 70 distributed points, the number of flip variables is also $O(n^2\ell)$. However, the number of flip
 71 variables is $\Theta(n^4\ell)$ if the points are in convex position (which is not the case for the challenge
 72 instances). Next, we describe the different types of clauses.

73 **Start.** For every edge variable $e(u, v, 0)$, we have the clause $e(u, v, 0)$ if $uv \in T_0$ and
 74 $\neg e(u, v, 0)$ if $uv \notin T_0$.

75 **Target.** For every edge variable $e(u, v, \ell)$, we have the clause $e(u, v, \ell)$ if $uv \in T_\ell$ and
 76 $\neg e(u, v, \ell)$ if $uv \notin T_\ell$.

77 **Flips need edges.** For every flip variable $f(u, v, u_2, v_2, i)$, we have the clause

$$78 \quad f(u, v, u_2, v_2, i) \implies e(u, v, i) \wedge e(u, v_2, i) \wedge e(u, u_2, i) \wedge e(v, v_2, i) \wedge e(v, u_2, i),$$

79 which easily translates to 5 binary CNF clauses.

80 **Flips keep edges.** For every flip variable $f(u, v, u_2, v_2, i)$, we have the clause

$$81 \quad f(u, v, u_2, v_2, i) \implies e(u_2, v_2, i+1) \wedge e(u, v_2, i+1) \wedge e(u, u_2, i+1) \wedge e(v, v_2, i+1) \wedge e(v, u_2, i+1),$$

82 which easily translates to 5 binary CNF clauses.

83 **Flips flip edges.** For every flip variable $f(u, v, u_2, v_2, i)$, we have the two clauses

$$84 \quad f(u, v, u_2, v_2, i) \implies e(u_2, v_2, i+1) \text{ and } f(u, v, u_2, v_2, i) \implies \neg e(u, v, i+1).$$

85 **Edge changes require flips.** The last type of clause is the only one that has more than 2
 86 variables in CNF form. It states that if the edge variable changes from triangulation i to
 87 $i+1$, then there must be a flip. The \vee below considers all values that produce valid flips.
 88 We have two such clauses for every edge variable:

$$89 \quad e(u, v, i) \wedge \neg e(u, v, i+1) \implies \bigvee_{u_2, v_2} f(u, v, u_2, v_2, i) \text{ and}$$

$$\neg e(u_2, v_2, i) \wedge e(u_2, v_2, i+1) \implies \bigvee_{u,v} f(u, v, u_2, v_2, i).$$

The number of variables and clauses grows very fast, even though the number of clauses is linear in the number of variables. Next, we show how to eliminate many variables from the model. All eliminated variables are assumed to be *false* in every clause, and the clauses that become tautologies are also eliminated. If a CNF clause becomes empty, then the problem is unsatisfiable.

Notice that we can easily eliminate $e(u, v, 0)$ for $uv \notin T_0$ and $e(u, v, \ell)$ for $uv \notin T_\ell$. This is, however, only a special case of a more general rule. The following theorem is easy to prove and implies that $\Omega(\log n)$ parallel flips are sometimes necessary to reconfigure two triangulations of n points, even when the points are in convex position. We say that two segments *cross* if they intersect at a point that is not an endpoint of either segment.

► **Theorem 1.** *Consider two triangulations T, T' of S such that a parallel flip transforms T into T' and a segment s with endpoints in S . Let χ, χ' respectively denote the number of edges of T, T' crossed by s . We then have $\chi' \geq \lfloor \chi/2 \rfloor$.*

Proof. Consider the sequence L of edges of T crossed by s in the order they cross the segment s . A parallel flip cannot remove two consecutive edges of L because they share a triangle, hence the theorem follows. ◀

Consequently, we only define the variable $e(u, v, i)$ when uv crosses strictly less than 2^i edges of T_0 and strictly less than $2^{\ell-i}$ target edges. We only define flip variables when a certain set of edge variables is defined. Namely, $f(u, v, u_2, v_2, i)$ is only defined when $uv, uv_2, uu_2, vv_2, vu_2$, are all defined at i and $u_2v_2, uv_2, uu_2, vv_2, vu_2$, are all defined at $i+1$.

2.2 Solution SAT Formulation

Next, we describe a SAT formulation for the following decision problem. Recall that an instance is a set S of n points and a list \mathcal{T} of input triangulations $\mathbf{T}_1, \dots, \mathbf{T}_{|\mathcal{T}|}$. The input of the decision problem is an instance and $|\mathcal{T}|$ integers $\ell_1, \dots, \ell_{|\mathcal{T}|}$. The output is whether there exists a solution $P_1, \dots, P_{|\mathcal{T}|}$ such that path P_i has length ℓ_i for all i .

We model the $|\mathcal{T}|$ paths $P_1, \dots, P_{|\mathcal{T}|}$ independently as before, starting path P_i at the input triangulation \mathbf{T}_i . The final triangulation of each path is unknown, but the same edge variables are used for the final triangulation of every path, since a valid solution requires that all paths end in the same triangulation. It is easy to see that the SAT formulation is satisfiable if and only if there exists a solution with the given lengths.

2.3 Lower Bound

In order to obtain an exact solution to an instance \mathcal{I} , we start by computing a lower bound to its objective value. We say that the *distance* between two triangulations T, T' is the length of the shortest path from T to T' . We create a complete directed graph $G(\mathcal{I})$ with edge weights as follows. The vertices are the triangulations \mathcal{T} and the weight of the edges are the distances between the corresponding triangulations. A *cycle packing* of G is a collection of vertex-disjoint directed cycles, i.e. a subset of edges such that each vertex has at most one outgoing and at most one incoming edge in the subset. The graph is directed to allow for cycles with only 2 edges. The *length of a cycle* is the sum of the lengths of its edges, and the *length of a cycle packing* is the sum of the lengths of its cycles. We have the following theorem.




Figure 3 (a) A cycle packing. (b) Illustration of the proof. In this example, $d_{1,2} \leq r_1 + r_2$, $d_{2,1} \leq r_2 + r_1$, $d_{3,4} \leq r_3 + r_4$, $d_{4,5} \leq r_4 + r_5$, and $d_{5,3} \leq r_5 + r_3$ by triangle inequality.

► **Theorem 2.** Given an instance \mathcal{I} , the objective value of a solution is at least the length of any cycle packing of $G(\mathcal{I})$ divided by 2.

Proof. Let $d_{i,j}$ denote the distance between the input triangulations $\mathbf{T}_i, \mathbf{T}_j$. Let C be a potential center and let r_i be the distance from C to \mathbf{T}_i . Consider a cycle $\mathbf{T}_1, \dots, \mathbf{T}_k$ in the cycle packing. By triangle inequality $d_{i,i+1} \leq r_i + r_{i+1}$ with indices taken modulo k (see Figure 3 (b)). Summing over the inequalities for i from 1 to k , we have that the length of the cycle is at most $2 \sum_i r_i$. Applying the same argument to every cycle, the theorem follows. \blacktriangleleft

143 2.4 The Exact Solver

First, we use the exact path formulation from Section 2.1 to calculate the distance between all $\binom{T}{2}$ pairs of input triangulations using a SAT solver (in our case, **CaDiCal**). It is easy to formulate the problem of finding a maximum length cycle packing as a weighted MaxSAT problem, which provides a lower bound b to the objective value (see Section 2.3). We solve this problem using a weighted MaxSAT solver (in our case **EvalMaxSAT**). We then use backtracking to list all integer solutions to $\ell_0, \dots, \ell_{|\mathcal{T}|} = b$ that satisfy $\ell_i + \ell_j \geq \text{distance}(T_i, T_j)$. We use the SAT formulation from Section 2.2 to test the existence of a solution with the given lengths $\ell_0, \dots, \ell_{|\mathcal{T}|}$, again using the **CaDiCal** SAT solver. If a solution is found, then it is optimal. Otherwise, we increment b and repeat. Notice that b is always a lower bound to the objective value. Hence, if a solution obtained by a heuristic attains this lower bound, then it is optimal.

155 **3** Heuristics

In this section, we describe different approaches that can be used when we do not need to guarantee the optimality of the solution. In Section 3.1, we present a conjecture that allows us to significantly increase the performance of the SAT solver. In Section 3.2, we show how to further increase the performance of the SAT solver using a heuristic coupled with some instance-independent preprocessed data. In Section 3.3, we show how to use the SAT solver to improve existing solutions. In Sections 3.4 and 3.5, we respectively show how to compute short paths and good solutions without a SAT solver.

163 3.1 Happy Edges Conjecture

¹⁶⁴ The happy edges conjecture [1] is a general conjecture that is *false* for some reconfiguration problems and *true* for others.

166 ▶ **Conjecture 3.** For any pair of configurations T, T' , there always exist a shortest path
167 between T, T' where the edges that are common to both T and T' appear in every intermediate
168 configuration.

169 The conjecture is *false* for triangulations under unit flips and arbitrary points [10] but
170 *true* when the points are in convex position [11]. Our experiments lead us to believe that the
171 conjecture is *true* for parallel flips and arbitrary point sets.

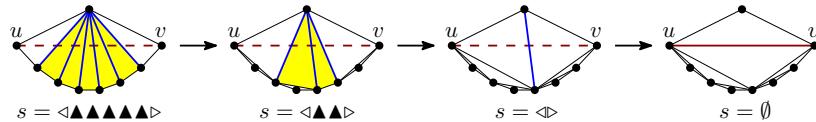
172 Enforcing that an edge never disappears and later reappears along a path actually makes
173 the SAT formulation harder to solve. However, there are some implications of the conjecture
174 that are very useful to make the SAT formulation shorter and easier to solve.

175 When computing a path of length ℓ from T_0 to T_ℓ using SAT, for every edge uv that
176 appears in both T_0 and T_ℓ , we add clauses $e(u, v, i)$ that force the variable to be true for all
177 i . More importantly, we then eliminate every edge variables corresponding to edges that
178 cross uv . The same idea can be applied to the SAT formulation that finds a solution, but
179 then only the edges that appear in all input triangulations are forced to be true for all i , and
180 again the edges that cross them are eliminated.

181 Furthermore, when computing a path, for every edge $uv \in T_\ell$, we eliminate flip variables
182 that remove uv , i.e. $f(u, v, u_2, v_2, i)$ for all u_2, v_2, i . Similarly, for every edge $u_2v_2 \in T_0$, we
183 eliminate flip variables that insert u_2v_2 , that is $f(u, v, u_2, v_2, i)$ for all u, v, i .

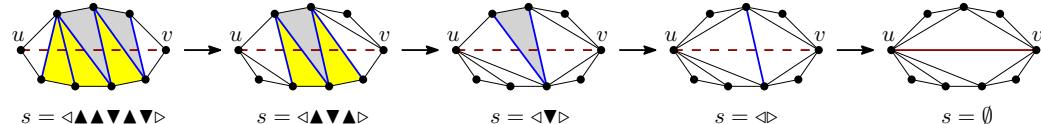
184 3.2 Crossing Lower Bound

185 Let $b(uv, T_0)$ denote the number of parallel flips needed to obtain an edge uv starting at a
186 triangulation T_0 . Clearly, when creating the SAT formulation for a path T_0, T_1, \dots we only
187 need to define the edge variable $e(u, v, i)$ for $i \geq b(T_0, uv)$. Theorem 1 implies that if an edge
188 uv crosses $\chi(uv, T_0)$ segments of T_0 , then $b(uv, T_0) \geq \lceil \log_2(\chi(uv, T_0) + 1) \rceil$. This bound is
189 tight when all the edges of T_0 that cross uv share a common endpoint and the endpoints are
190 in convex position (Figure 4). However, there are different ways in which the edges of T_0
191 may cross uv that may imply higher values of $b(uv, T_0)$.



192 ■ **Figure 4** A path of length 3 to insert an edge that had 6 crossings.

193 We consider estimations of $b(uv, T_0)$ based on the sequence of triangles that contain
194 an upper or a lower edge (Figure 5). An edge uv that crosses $\chi(uv, T_0)$ segments of T_0 is
195 translated into a *string* $s = s(uv, T_0)$ of $\chi(uv, T_0) + 1$ symbols in the alphabet $\{\nabla, \blacktriangle, \triangleleft, \triangleright\}$,
196 according to which sides of uv contain the edges that are not crossed by uv in each triangle.



197 ■ **Figure 5** A path of length 4 to insert an edge that had 6 crossings. Each triangle is labeled and
198 colored as containing an upper or a lower edge.

199 The unit flips on T_0 have equivalent *productions* on s that replace the substring on the
200 left-hand side by the substring on the right-hand side. We define a *substring* as a contiguous

201 subsequence. Flipping the two extreme triangles translates to the *extreme productions*
 202 $\triangleleft \blacktriangle \rightarrow \triangleleft$, $\triangleleft \blacktriangledown \rightarrow \triangleleft$, $\blacktriangle \triangleright \rightarrow \triangleright$, and $\blacktriangledown \triangleright \rightarrow \triangleright$. Flipping intermediate triangles with the same
 203 orientation translates to $\blacktriangle \blacktriangle \rightarrow \blacktriangle$ and $\blacktriangledown \blacktriangledown \rightarrow \blacktriangledown$, while flipping intermediate triangles
 204 of different orientations translates to $\blacktriangledown \blacktriangle \rightarrow \blacktriangle \blacktriangledown$ and $\blacktriangle \blacktriangledown \rightarrow \blacktriangledown \blacktriangle$. The last flip to insert
 205 the edge uv is $\triangleleft \triangleright \rightarrow \emptyset$. There are other productions that may increase the string length
 206 such as $\blacktriangle \rightarrow \blacktriangle \blacktriangle$, which consist of exchanging the left-hand and right-hand side of some
 207 aforementioned productions, but we show that we can always obtain shortest paths without
 208 using them.

209 A parallel flip means that we may apply several productions simultaneously as long as
 210 their left-hand side corresponds to disjoint substrings. We call the application of several
 211 productions a *rewriting*. A *rewriting sequence* of length ℓ is a sequence of $\ell + 1$ strings
 212 connected by rewriting operations and ending at the empty string \emptyset . We want to find shortest
 213 rewriting sequences. Let $b(s)$ be the length of the shortest rewriting sequence of the string s .

214 Notice that if the points are not in convex position, then some productions may correspond
 215 to invalid flips. Hence, assuming convex position provides a lower bound $b(s(uv, T_0)) \leq$
 216 $b(uv, T_0)$. Furthermore, the productions assume that there is an unbounded number of points,
 217 so that new triangles may be created freely.

218 We use a heuristic together with a depth first search exact solution to estimate $b(s)$. The
 219 heuristic is very involved and sometimes gives incorrect bounds (50182 wrong values for
 220 33554432 tested strings, which is roughly 1/600 wrong bounds, the smallest wrong bound
 221 being for $\triangleleft \blacktriangledown \blacktriangle \blacktriangle \blacktriangledown \blacktriangle \blacktriangle \blacktriangle \triangleright$). To fix some wrong bounds, we use an instance-independent
 222 preprocessing to compute tight values of $b(s)$ for small enough strings s and store only the
 223 values where the heuristic incorrectly computes $b(s)$ in a file. This file is loaded by our solver
 224 and stored in a hash table.

225 3.3 Improving a Solution

226 Given a solution $P_1, \dots, P_{|\mathcal{T}|}$, we may improve it as follows. We choose a random path P_i
 227 and use the SAT formulation to find a solution where the length of P_i is decremented and all
 228 other lengths remain the same. We repeat this as often as necessary. Notice that the method
 229 may converge to a locally optimal minimal list of lengths that is not globally optimal.

230 If the number of clauses is too large, there are two different approaches that we can take
 231 (and they may be combined). We may force the new solution to be close to the original one
 232 by only creating edge variables that cross few edges in the corresponding triangulation of the
 233 previous solution.

234 We may also trim the solution to a certain radius r , by only rebuilding the portion of the
 235 solution that is within r steps from the center. In this case, it is helpful to use MaxSAT first,
 236 in order to reduce the number of unit-flips performed in the last steps as follows. Given a
 237 path T_0, \dots, T_ℓ , we first use a MaxSAT solver to find the path of length ℓ that minimizes
 238 the number of unit flips performed from $T_{\ell-1}$ to T_ℓ . The MaxSAT formulation is equal to
 239 the SAT formulation with soft clauses $\neg f(u, v, u_2, v_2, \ell)$ for each last step flip variable. We
 240 then find the path of length $\ell - 1$ from T_0 to the $T_{\ell-1}$ of the previous path that minimizes
 241 the number of unit flips performed from $T_{\ell-2}$ to $T_{\ell-1}$. We continue this way for r steps.

242 3.4 Path Heuristic

243 The SAT formulation finds the shortest path connecting two triangulations reasonably fast,
 244 but it may be too slow for some usages. We also designed a heuristic that produces reasonably

245 short paths quickly. We are given two triangulations T_0, T' and the goal is to find a short
 246 path from T_0 to T' .

247 We use a greedy approach that iteratively obtains a triangulation T_{i+1} from a triangulation
 248 T_i as follows. Let F denote the set of possible unit flips in T_i . More precisely, the elements
 249 of F are pairs e, e' of edges such that a unit flip from T_i removes e and inserts e' . We then
 250 build a graph $G(F)$ with vertex set F and edges between two unit flips that share a triangle.
 251 Notice that the possible parallel flips correspond to independent sets in $G(F)$. We assign
 252 a weight to the vertices as follows. The weight of a vertex e, e' is the number of edges in
 253 T' crossed by e' minus the number of edges in T' crossed by e . Vertices of zero or negative
 254 weight are eliminated.

255 We then proceed to greedily find an independent set I in $G(F)$. We iteratively add to I
 256 the *unmarked* vertex of maximum weight, breaking ties by minimum degree. We then *mark*
 257 all the vertices in the closed neighborhood of I . We repeat until all vertices are marked.

258 This iterative approach is repeated until we reach T' , which will happen in $O(n^2)$ flips
 259 because of the following Theorem from [7].

260 ► **Theorem 4.** *Let T, T' be two triangulations of the same point set. If $T \neq T'$, then there
 261 exists a unit flip that replaces an edge $e \in T$ by an edge e' such that e crosses more edges of
 262 T' than e' does.*

263 We further improve the heuristic using the squeaky wheel paradigm [8]. Initially, we
 264 assign weight 1 to every edge. We modify the definition of the weight of a flip as follows.
 265 The weight of a flip e, e' is the sum of the weights of the edges in T' crossed by e' minus the
 266 sum of the weights of the edges in T' crossed by e . When we reach triangulation $T_{i+1} = T'$
 267 from a triangulation T_i we increment the weight of all edges in T' that are not in T_{i+1} . We
 268 repeat the greedy algorithm with the new weights, keeping the best solution found. This
 269 procedure is repeated multiple times.

270 3.5 Solution Heuristic

271 We use the following strategy to obtain reasonably good initial solutions that we may use
 272 as a basis to improve with the aforementioned methods. We start by adding the Delaunay
 273 triangulation as a candidate center. We then build other candidate centers by performing
 274 flips starting from the Delaunay triangulation. Given an edge e and a triangulation T , let
 275 $\chi(e, T)$ denote the number of segments of T crossed by e . We repeatedly perform unit flips
 276 that remove an edge e and add an edge e' maximizing

$$278 \quad \sum_{T \in \mathcal{T}} \chi(e)^p - \chi(e')^p,$$

279 as long as the value of the sum is positive. We add the triangulation we obtained to the set
 280 of candidate center and repeat the process for a different value of p .

281 We calculate the paths from the candidate centers to each solution, building a solution
 282 pool. For the next step it will be useful to have a small number of unit-flips in the last flip
 283 of the path from the input triangulation to the center. To do that, we either use the greedy
 284 heuristic or a MaxSAT formulation.

285 To improve the solution pool, we pick a solution from the pool and look at the triangulations
 286 that are one flip away from the center in each path. For each such triangulation, we compute
 287 the distance to the input triangulations and add the solution to the pool as before.

288 4 Results

289 In this section, we present the computational results that we obtained with our implementation
 290 of the aforementioned algorithms and heuristics. In Section 4.1, we present the results on
 291 computing short paths between two given triangulations. In Section 4.2, we present our
 292 exact solver, with and without the happy edges conjecture. In Section 4.3, we present the
 293 heuristics used to find solutions to the instances that we could not solve exactly.

294 The solvers were coded in C++ and compiled with GCC and run a single thread.
 295 During the competition, they were executed on several Linux computers, either using GNU
 296 Parallel [12] for local executions or Slurm [13] for cluster executions. It was very useful to
 297 have access to machines with 128GB or more RAM to be able to solve large SAT formulations
 298 with CaDiCal [3], which has been able to solve SAT instances with more than 50 million
 299 variables and 500 million clauses. The time measurements on this paper have all been taken
 300 on an AMD Ryzen 9 9900X CPU and ASUS TUF GAMING B650M motherboard with
 301 128GB of RAM running Fedora Core 43.

302 4.1 Path Calculation

303 Computing short paths between two given triangulations is a key component to obtain good
 304 solutions. Typically, these paths are computed with an input triangulation as one extreme,
 305 and a triangulation that makes a reasonably good center as the other extreme. In this section,
 306 we use the Delaunay triangulation as one extreme, because it is a well defined triangulation
 307 that makes a reasonably good (but not very good) center. Table 1 shows the length of the
 308 path and the running time of different heuristics and SAT solutions with and without the
 309 happy edges conjecture. The paths are calculated from the Delaunay triangulation to the
 310 first input triangulation of several instances. The SAT paths are obtained by first running
 311 the squeaky wheel heuristic in both directions, and then iteratively decreasing the path
 312 length. Notice that the running time of the heuristics is much smaller, while the result is
 313 rarely more than 1 unit away from the optimal distance, especially if we take the minimum
 314 of both directions. We run the squeaky wheel heuristic for at most 16 iterations, but stops
 315 earlier if the length increases from one iteration to the next.

334 4.2 Exact Solutions

335 This section presents the exact solver that we used to solve most instances exactly. Figure 6
 336 shows the number of exact solutions found as a function of the running time, both without
 337 and with the happy edges conjecture (the value of the solutions found in both cases is the
 338 same, hence the conjecture stands). Notice that the impact of the happy edges conjecture is
 339 less significant than when computing only paths, as there are few common edges on all input
 340 triangulations. We remark that during the competition, we managed to find exact solutions
 341 to 189 of the 201 small instances without assuming the happy edges conjecture.

344 4.3 Heuristic Solutions

345 The typical process to solve instances that we did not solve exactly consists of several steps.
 346 First, we use the heuristic from Section 3.5 to obtain a reasonably good center. Notice that
 347 no SAT solver is used in this part. The evolution of the objective value for this step is shown
 348 in Figure 7. The name of the `rirs` instances is composed of two values. The first is the
 349 number of points and the second is the number of input triangulations.

n	Greedy forward		Greedy backward		Squeaky forward		Squeaky backward		SAT happy		SAT exact	
	ℓ	t	ℓ	t	ℓ	t	ℓ	t	ℓ	t	ℓ	t
500	13	4.2	12	3.7	13	45	12	18	12	1512	12	10191
1000	12	7.8	12	6.8	12	25	12	38	11	6806	11	78264
1500	13	13	12	12	12	52	12	23	11	19302	11	201353
2000	15	21	13	18	13	65	13	104	13	16937	13	485177
3000	15	36	15	33	14	193	15	433	14	140729	.	.
4000	18	57	16	51	16	577	15	655	15	101402	.	.
5000	18	70	17	70	16	238	17	278	16	1037840	.	.
6000	17	93	16	83	17	1071	16	644	16	271081	.	.
7000	17	94	17	96	17	180	17	407	16	471389	.	.

Table 1 Length and computation time (in milliseconds) from the Delaunay triangulation to the first input triangulation of the `rirs-n-20` instance for different values of the number of points n . The columns respectively correspond to the greedy heuristic forward and backward, the squeaky wheel heuristic forward and backward, the SAT solution with the happy edges conjecture, and the SAT solution without the happy edges conjecture, unless it takes too long. The best lengths found are shown in bold.


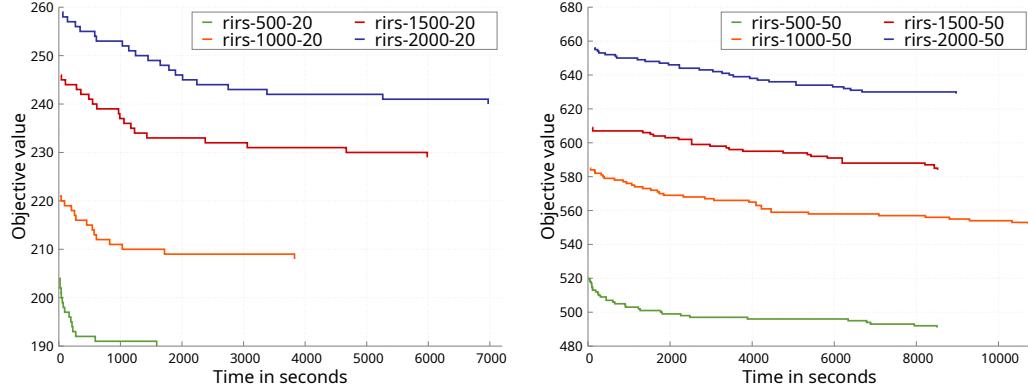


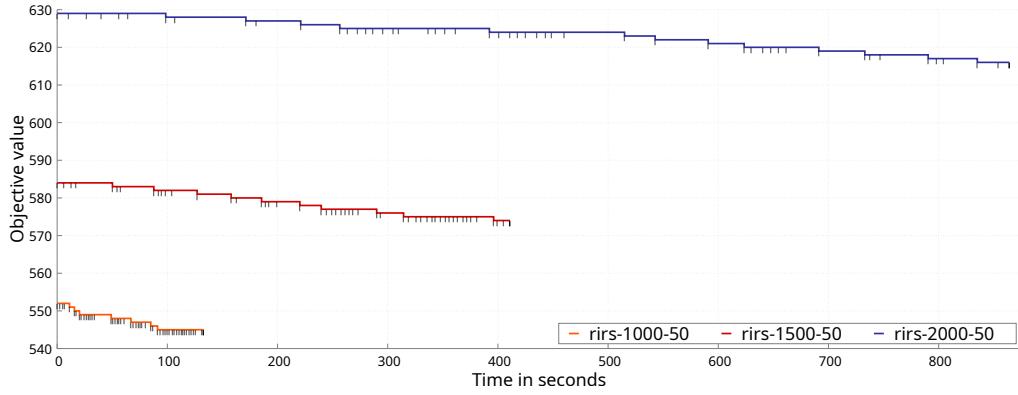
Figure 6 Number of exact solutions found as a function of the running time over 3 hours of execution with and without assuming the happy edges conjecture.

Second, we build a new solution that keeps the same center but we recalculate the paths using the SAT formulation from 2.1, assuming the happy edges conjecture (Section 3.1) and inexact lower bounds (Section 3.2). This will typically reduce the objective value by 1 to 4 percent. The calculation of 50 paths in each solution is shown in Figure 8. Notice that the running time increases rapidly with the number of points and that finding a shorter path (satisfiable SAT problem) is typically slower than when no shorter path exists (unsatisfiable SAT problems).

Third, we improve the solution using the SAT formulation from Section 2.2, assuming the happy edges conjecture (Section 3.1) and inexact lower bounds (Section 3.2). This part requires adjusting a large number of parameters in order to obtain SAT problems that are not too hard. The parameters include the distance to the previous center, the distance to the previous path, and whether the solution will be trimmed. Optionally but recommended when the problem is trimmed, we use MaxSAT to reduce the number of unit flips close to the center. The improvement of some solutions over time is shown in Figure 9. Notice that



350 ■ **Figure 7** Evolution of the best solution found by the heuristic solver without any SAT solver for
351 different instances.




389 ■ **Figure 8** Path by path improvement of the heuristic solutions using a SAT solver, where each path
390 is recalculated but the center is kept unchanged. Each black vertical bar represents the beginning of
391 the computation of a path.

399 there is a significant preprocessing time to calculate edge variables with a limited number of
400 crossings. This preprocessing is applied initially and after every improvement. Also notice
401 that unsatisfiable SAT problems, which mean no improvement in the solution, are solved
402 much faster than satisfiable ones.

379 5 Conclusion and Open Problems

380 We were surprised that we managed to solve so many instances exactly and how well an
381 heuristic approximates the exact distance. We believe the following factors help explain the
382 strong practical performance:

- 383 ■ The short path length that allows a somewhat small number of variables.
- 384 ■ A SAT model where most clauses have size 2.
- 385 ■ The ability to eliminate many variables through several arguments.
- 386 ■ The fact that the happy edges conjecture is either true or at least holds in most practical
387 cases.



373 ■ **Figure 9** Improvement of the whole solution using a SAT solver and reducing the length of a
 374 random path by one unit at a time. We constrain the edges in the new solution to cross at most
 375 3 edges of the corresponding triangulation in the solution (except for the instance **rirs-500-20**,
 376 where the slower but more effective execution with the parameter set to at most 7 edges is also
 377 pictured) but do not use trimming. Each black vertical bar represents a new path length that we
 378 try to decrement.

388 The short path lengths are somewhat surprising, given that an $\Omega(n)$ lower bound is
 389 presented in [6], in contrast to the $\Theta(\log n)$ bound for combinatorial triangulations [9], where
 390 we are allowed to flip non-convex quadrilaterals. In the challenge instances, we observed
 391 path lengths that are roughly $O(\log n)$.

392 Still, there are **random** instances with only 160 points and 20 input triangulations and
 393 **woc** instances with only 185 points and 6 triangulations that we could not solve exactly. Also,
 394 we still managed to improve solutions to instances with as few as 320 points and 20 input
 395 triangulations months after the beginning of the challenge.

396 The challenge instances did not have many points in convex position. Surprisingly, the
 397 case where all points are in convex position is the hardest for our SAT formulation, as there
 398 are $\Theta(n^4)$ empty convex quadrilaterals. We wonder if a different model works better when
 399 all and also when most points are in convex position.

400 We believe that the same problem with unit flips is significantly harder because the long
 401 path lengths make the SAT formulation much more complex and removing a happy edge can
 402 reduce a path length from $\Theta(n^2)$ to $\Theta(n)$.

403 Many theoretical open problems remain, such as proving Conjecture 3. We also wonder
 404 if parallel flip distance problem is NP-hard in general and in convex position, as well as

405 the unit flip distance problem in the convex case (the problem is NP-hard for general point
 406 sets [10]). We also do not know if the problem of calculating $b(s)$ (see Section 3.2) can be
 407 solved in polynomial time, possibly using dynamic programming.

 408 **References**

- 409 1 Oswin Aichholzer, Brad Ballinger, Therese Biedl, Mirela Damian, Erik D Demaine, Matias
 410 Korman, Anna Lubiw, Jayson Lynch, Josef Tkadlec, and Yushi Uno. Reconfiguration of
 411 non-crossing spanning trees, 2022. [arXiv:2206.03879](https://arxiv.org/abs/2206.03879).
- 412 2 Florent Avellaneda. A short description of the solver EvalMaxSAT. *MaxSAT Evaluation*,
 413 8:364, 2020.
- 414 3 Armin Biere, Tobias Faller, Katalin Fazekas, Mathias Fleury, Nils Froleyks, and Florian Pollitt.
 415 CaDiCaL 2.0. In *Computer Aided Verification - 36th International Conference, CAV*, volume
 416 14681 of *LNCS*, pages 133–152, 2024. doi:[10.1007/978-3-031-65627-9_7](https://doi.org/10.1007/978-3-031-65627-9_7).
- 417 4 Ruy Fabila-Monroy, Clemens Huemer, and Dieter Mitsche. The number of empty four-
 418 gons in random point sets. *Electronic Notes in Discrete Mathematics*, 46:161–168, 2014.
 419 doi:[10.1016/j.endm.2014.08.022](https://doi.org/10.1016/j.endm.2014.08.022).
- 420 5 Sándor P. Fekete, Phillip Keldenich, Dominik Krupke, and Stefan Schirra. Minimum coverage
 421 by convex polygons: The CG:SHOP challenge 2023, 2023. [arXiv:2303.07007](https://arxiv.org/abs/2303.07007).
- 422 6 Jérôme Galtier, Ferran Hurtado, Marc Noy, Stéphane Pérennes, and Jorge Urrutia.
 423 Simultaneous edge flipping in triangulations. *International Journal of Computational Geometry
 & Applications*, 13(02):113–133, 2003.
- 424 7 Sabine Hanke, Thomas Ottmann, and Sven Schuierer. The edge-flipping distance of
 425 triangulations. *Journal of Universal Computer Science*, 2(8):570–579, 1996.
- 426 8 David E. Joslin and David P. Clements. Squeaky wheel optimization. *Journal of Artificial
 427 Intelligence Research*, 10:353–373, 1999.
- 428 9 Tanvir Kaykobad. An interlaced algorithm for transforming plane triangulations using
 429 simultaneous flips. *Journal of Computational Geometry*, 16(1):517–550, 2025.
- 430 10 Alexander Pilz. Flip distance between triangulations of a planar point set is APX-hard.
 431 *Computational Geometry*, 47(5):589–604, 2014. doi:[10.1016/j.comgeo.2014.01.001](https://doi.org/10.1016/j.comgeo.2014.01.001).
- 432 11 Daniel D. Sleator and William P. Tarjan, Robert E. and Thurston. Rotation distance,
 433 triangulations, and hyperbolic geometry. In *Proceedings of the eighteenth annual ACM
 434 symposium on Theory of computing*, pages 122–135, 1986.
- 435 12 O. Tange. Gnu parallel - the command-line power tool. *:login: The USENIX Magazine*,
 436 36(1):42–47, Feb 2011. URL: <http://www.gnu.org/s/parallel>.
- 437 13 Andy B Yoo, Morris A Jette, and Mark Grondona. Slurm: Simple linux utility for resource
 438 management. In *Workshop on job scheduling strategies for parallel processing*, pages 44–60,
 439 2003.
- 440