

Shadoks Approach to Parallel Reconfiguration of Triangulations

Guilherme D. da Fonseca  

LIS, Aix-Marseille Université

Fabien Feschet  

LIMOS, Université Clermont Auvergne

Yan Gerard  

LIMOS, Université Clermont Auvergne

1 Abstract

We describe the heuristics used by the Shadoks team in the CGSHOP 2026 Challenge. The Challenge consists of 250 instances, each being a list of triangulations of the same point set. The goal is to find short paths from a chosen central triangulation to each instance triangulation, where a parallel flip operation connects consecutive triangulations in each path. We use several heuristics and exact methods to solve the problem, based on greedy approaches as well as SAT and MaxSAT formulations.

2012 ACM Subject Classification Theory of computation → Computational geometry

Keywords and phrases Heuristic, exact algorithm, SAT solver, MaxSAT solver, computational geometry

Category CG Challenge


Related Version arxiv.org/abs/

Supplementary Material (Source Code): <https://github.com/gfonsecabr/shadoks-CGSHOP2026>

Acknowledgements We would like to thank the Challenge organizers and other competitors for their time, feedback, and making this whole event possible. We would like to thank Hélène Toussaint, Raphaël Amato, Boris Lonjon, and William Guyot-Lénat from LIMOS, as well as the Qarma and TALEP teams and Manuel Bertrand from LIS, who continue to make the computational resources of the LIMOS and LIS clusters available to our research. We would also like to thank Aldo Gonzalez-Lorenzo for the very useful discussion on SAT models.

8 1 Introduction

In this paper, we outline the exact methods and heuristics that we employed. We start with some definitions that allow us to describe the problem. Throughout, we consider



9  **Figure 1** Our best centers to instances `random_78_40_10`, `woc-70-random-9a7d18d3`,
10 `woc-90-tsplib`, and `rirs-500-50-23d00ec5`, respectively.

13 triangulations of the same point set $P \subset \mathbb{R}^2$. Given a triangulation T , a *unit flip* is the
 14 operation that removes an edge $e \in T$ and adds an edge e' , obtaining a new triangulation
 15 $T' = T \setminus \{e\} \cup \{e'\}$. Similarly, a *parallel flip* removes a set of edges $E \subset T$ and adds a set of
 16 edges E' , in a way that $T' = T \setminus E \cup E'$ is a triangulation, with the condition that no two
 17 edges of E are in the same triangle in T . A *path* of length ℓ is a sequence of triangulations
 18 T_0, \dots, T_ℓ such that for all i , the triangulation T_{i+1} is obtained from T_i by performing a
 19 parallel flip.

20 An *instance* consists of a set $P \subset \mathbb{R}^2$ of n points and a set of triangulations $\mathcal{T} = T_1, \dots, T_{|\mathcal{T}|}$
 21 , called *input triangulations*. A *solution* is a set of paths $P_1, \dots, P_{|\mathcal{T}|}$ such that P_i starts at
 22 T_i for all i and all paths end in a common triangulation called *center*. The goal is to find a
 23 solution that minimizes the *objective value* defined as the sum of the lengths of its paths.

24 During the competition, the organizers provided a total of 250 instances, with P ranging
 25 from 15 to 12,500 points and \mathcal{T} ranging from 2 to 200 triangulations. The 250 instances are
 26 divided into three classes: 100 **random** instances, 101 **woc** instances, and 49 **rirr** instances.
 27 The former two instances have up to 320 points and 2 to 20 input triangulations (hence,
 28 we call them **small** instances), while the latter have 500 to 12500 points and 20 to 200
 29 input triangulations. The centers of some of our best solutions are presented in Figure 1.
 30 Additional details about the challenge can be found in the organizers' survey paper [4].

31 Our best solvers heavily rely on the SAT solver **CaDiCaL** [5] and the MaxSAT solver
 32 **EvalMaxSAT** [2]. Nevertheless, we also developed heuristics that do not rely on any external
 33 solver, which are important to find initial solutions to some large instances, which are then
 34 improved by roughly 10% using SAT and MaxSAT solvers. Furthermore, we managed to
 35 solve 189 of the 201 **small** instances exactly by repeatedly using the SAT solver as well as
 36 some lower bounds.

37 Mention the other teams strategy here...

38 We describe our exact algorithms in Section 2, the heuristics in Section 3, and discuss
 39 the results we obtained in Section 4. Concluding remarks and open problems are presented
 40 in Section 5.

41 2 Exact Algorithms


42 This section describes all elements of our exact solver, many of which are also used in the
 43 heuristic solvers. We first show how to use a SAT solver to compute shortest paths between
 44 two triangulations (Section 2.1). We then show how to extend this result to test if a solution
 45 with a list of path lengths exists (Section 2.2). We show how to obtain lower bounds in
 46 Section 2.3 and put the previous elements together to describe our exact solver in Section 2.4.

47 2.1 Path SAT Formulation

48 We now describe a SAT formulation for the following decision problem. The input is a set P
 49 of n points, an integer ℓ and two triangulations T_0, T_ℓ . The output is whether there exists a
 50 path T_0, \dots, T_ℓ of length ℓ .

51 We define two types of variables. For $i = 0, \dots, \ell$ and for $u \neq v \in P$, we define an *edge*
 52 *variable* $e(u, v, i)$. The variable $e(u, v, i)$ represents that the edge uv is in the triangulation T_i .
 53 There are $O(n^2\ell)$ edge variables. It would be possible to define a SAT formulation using only
 54 such variables. However, a SAT formulation that performed much better in our experiments
 55 uses a second type of variable.

56 We say that a convex quadrilateral is *empty* if it contains no point of P except for its
 57 vertices. For $i = 0, \dots, \ell$ and for $u \neq v \neq u_2 \neq v_2 \in P$ such that u, u_2, v, v_2 form an empty



56 **Figure 2** Illustration of a flip variable $f(u, v, u_2, v_2, i)$.

59 convex quadrilateral, we define a *flip variable* $f(u, v, u_2, v_2, i)$. The variable $f(u, v, u_2, v_2, i)$
60 represents a unit flip such that the edge uv is in triangulation T_i and u_2v_2 is in triangulation
61 T_{i+1} , as shown in Figure 2. Notice that if the points are uniformly distributed, then
62 the number of empty convex quadrilaterals is $\Theta(n^2)$ [3], which means that for uniformly
63 distributed points, the number of flip variables is also $O(n^2\ell)$. However, the number of flip
64 variables is $\Theta(n^4\ell)$ if the points are in convex position (which is not the case for the challenge
65 instances). Next, we describe the different types of clauses.

66 **Start.** For every edge variable $e(u, v, 0)$, we have the clause $e(u, v, 0)$ if $uv \in T_0$ and
67 $\neg e(u, v, 0)$ if $uv \notin T_0$.

68 **Target.** For every edge variable $e(u, v, \ell)$, we have the clause $e(u, v, \ell)$ if $uv \in T_\ell$ and
69 $\neg e(u, v, \ell)$ if $uv \notin T_\ell$.

70 **Flips need edges.** For every flip variable $f(u, v, u_2, v_2, i)$, we have the clause

$$71 \quad f(u, v, u_2, v_2, i) \implies e(u, v, i) \wedge e(u, v_2, i) \wedge e(u, u_2, i) \wedge e(v, v_2, i) \wedge e(v, u_2, i),$$

72 which easily translates to 5 binary CNF clauses.

73 **Flips keep edges.** For every flip variable $f(u, v, u_2, v_2, i)$, we have the clause

$$74 \quad f(u, v, u_2, v_2, i) \implies e(u_2, v_2, i+1) \wedge e(u, v_2, i+1) \wedge e(u, u_2, i+1) \wedge e(v, v_2, i+1) \wedge e(v, u_2, i+1),$$

75 which easily translates to 5 binary CNF clauses.

76 **Flips flip edges.** For every flip variable $f(u, v, u_2, v_2, i)$, we have the two clauses

$$77 \quad f(u, v, u_2, v_2, i) \implies e(u_2, v_2, i+1) \text{ and } f(u, v, u_2, v_2, i) \implies \neg e(u, v, i+1).$$

78 **Edge changes require flips.** The last type of clause is the only one that has more than 2
79 variables in CNF form. It states that if the edge variable changes from triangulation i to
80 $i+1$, then there must be a flip. The \bigvee below considers all values that produce valid flips.
81 We have two such clauses for every edge variable:

$$82 \quad e(u, v, i) \wedge \neg e(u, v, i+1) \implies \bigvee_{u_2, v_2} f(u, v, u_2, v_2, i) \text{ and}$$

$$83 \quad \neg e(u_2, v_2, i) \wedge e(u_2, v_2, i+1) \implies \bigvee_{u, v} f(u, v, u_2, v_2, i).$$

85 The number of variables and clauses grows very fast, even though the number of clauses
86 is linear in the number of variables. Next, we show how to eliminate many variables from the
87 model. All eliminated variables are assumed to be *false* in every clause, and the clauses that

88 become tautologies are also eliminated. If a CNF clause becomes empty, then the problem is
 89 unsatisfiable.

90 Notice that we can easily eliminate $e(u, v, 0)$ for $uv \notin T_0$ and $e(u, v, \ell)$ for $uv \notin T_\ell$. This
 91 is, however, only a special case of a more general rule. The following theorem is easy to
 92 prove and implies that $\Omega(\log n)$ parallel flips are sometimes necessary to reconfigure two
 93 triangulations of n points, even when the points are in convex position. We say that two
 94 segments *cross* if they intersect at a point that is not an endpoint of either segment.

95 ► **Theorem 1.** *Consider two triangulations T, T' of P such that a parallel flip transforms T
 96 into T' and a segment s with endpoints in P . Let χ, χ' respectively denote the number of
 97 edges of T, T' crossed by s . We then have $\chi' \geq \lfloor \chi/2 \rfloor$.*

98 **Proof.** Consider the list L of edges of T crossed by s in the order they cross the segment s .
 99 A parallel flip cannot remove two consecutive edges of L because they share a triangle, hence
 100 the theorem follows. ◀

101 Consequently, we only define the variable $e(u, v, i)$ when uv crosses strictly less than
 102 2^i edges of T_0 and strictly less than $2^{\ell-i}$ target edges. We only define flip variables when
 103 a certain set of edge variables is defined. Namely, $f(u, v, u_2, v_2, i)$ is only defined when
 104 $uv, uv_2, uu_2, vv_2, vu_2$, are all defined at i and $u_2v_2, uv_2, uu_2, vv_2, vu_2$, are all defined at $i+1$.

105 2.2 Solution SAT Formulation

106 Next, we describe a SAT formulation for the following decision problem. Recall that an
 107 instance is a set P of n points and a list \mathcal{T} of triangulations $T_0, \dots, T_{|\mathcal{T}|}$. The input of the
 108 decision problem is an instance and $|\mathcal{T}|$ integers $\ell_0, \dots, \ell_{|\mathcal{T}|}$. The output is whether there
 109 exists a solution $P_0, \dots, P_{|\mathcal{T}|}$ such that path P_i has length ℓ_i for all i .

110 We model the $|\mathcal{T}|$ paths $P_0, \dots, P_{|\mathcal{T}|}$ independently as before, starting path P_i at the
 111 input triangulation T_i . The final triangulation of each path is unknown, but the same edge
 112 variables are used for the final triangulation of every path, since a valid solution requires
 113 that all paths end in the same triangulation. It is easy to see that the SAT formulation is
 114 satisfiable if and only if there exists a solution with the given lengths.

115 2.3 Lower Bound

116 In order to obtain an exact solution to an instance \mathcal{I} , we start by computing a lower bound
 117 to its objective value. We say that the *distance* between two triangulations T, T' is the length
 118 of the shortest path from T to T' . We create a complete directed graph $G(\mathcal{I})$ with edge
 119 weights as follows. The vertices are the triangulations \mathcal{T} and the weight of the edges are the
 120 distances between the corresponding triangulations. A *cycle packing* of G is a collection of
 121 vertex-disjoint directed cycles, i.e. a subset of edges such that each vertex has at most one
 122 outgoing and at most one incoming edge in the subset. The graph is directed to allow for
 123 cycles with only 2 edges. The *length of a cycle* is the sum of the lengths of its edges, and
 124 the *length of a cycle packing* is the sum of the lengths of its cycles. We have the following
 125 theorem.

128 ► **Theorem 2.** *Given an instance \mathcal{I} , the objective value of a solution is at least the length of
 129 any cycle packing of $G(\mathcal{I})$ divided by 2.*

130 **Proof.** Let $d_{i,j}$ denote the distance between triangulations T_i, T_j . Let C be a potential center
 131 and let r_i be the distance from C to T_i . Consider a cycle T_1, \dots, T_k in the cycle packing.




Figure 3 (a) A cycle packing. (b) Illustration of the proof. In this example, $d_{1,2} \leq r_1 + r_2$, $d_{2,1} \leq r_2 + r_1$, $d_{3,4} \leq r_3 + r_4$, $d_{4,5} \leq r_4 + r_5$, and $d_{5,3} \leq r_5 + r_3$ by triangle inequality.

¹³² By triangle inequality $d_{i,i+1} \leq r_i + r_{i+1}$ with indices taken modulo k (see Figure 3 (b)).
¹³³ Summing over the inequalities for i from 1 to k , we have that the length of the cycle is at
¹³⁴ most $2 \sum_i r_i$. Applying the same argument to every cycle, the theorem follows. ◀

135 2.4 The Exact Solver

First, we use the exact path formulation from Section 2.1 to calculate the distance between all $\binom{T}{2}$ input triangulations using a SAT solver (in our case, **CaDiCa1**). It is easy to formulate the problem of finding a maximum length cycle packing as a weighted MaxSAT problem, which provides a lower bound b to the objective value (see Section 2.3). We solve this problem using a weighted MaxSAT solver (in our case **EvalMaxSAT**). We then use backtracking to list all integer solutions to $\ell_0, \dots, \ell_{|\mathcal{T}|} = b$ that satisfy $\ell_i + \ell_j \geq \text{distance}(T_i, T_j)$. We use the SAT formulation from Section 2.2 to test the existence of a solution with the given lengths $\ell_0, \dots, \ell_{|\mathcal{T}|}$, again using the **CaDiCa1** SAT solver. If a solution is found, then it is optimal. Otherwise, we increment b and repeat. Notice that b is always a lower bound to the objective value. Hence, if a solution obtained by a heuristic attains this lower bound, then it is optimal.

146 **3** Heuristics

In this Section, we describe different approaches that can be used when we do not need to guarantee the optimality of the solution. In Section 3.1, we present a conjecture that allows us to significantly increase the performance of the SAT solver. In Section 3.2, we show how to further increase the performance of the SAT solver using a heuristic coupled with some instance-independent preprocessed data. In Section 3.3, we show how to use the SAT solver to improve existing solutions. In Sections 3.4 and 3.5, we respectively show how to compute short paths and good solutions without a SAT solver.

154 3.1 Happy Edges Conjecture

¹⁵⁵ The happy edges conjecture [1] is a general conjecture that is *false* for some reconfiguration problems and *true* for others.

► **Conjecture 3.** For any pair of configurations T, T' , there always exist a shortest path between T, T' where the edges that are common to both T and T' appear in every intermediate configuration.

The conjecture is *false* for triangulations under unit flips and arbitrary points [10] but *true* when the points are in convex position [11]. Our experiments lead us to believe that the conjecture is *true* for parallel flips and arbitrary point sets.

Our SAT formulation does not become easier to solve if we force the condition that edges are not allowed to disappear and subsequently reappear on a path. However, there are some implications of the conjecture that are very useful to make the SAT formulation shorter and easier to solve.

When computing a path of length ℓ from T_0 to T_ℓ using SAT, for every edge uv that appears in both T_0 and T_ℓ , we add clauses $e(u, v, i)$ that force the variable to be true for all i . More importantly, we then eliminate every edge variables corresponding to edges that cross uv . The same idea can be applied to the SAT formulation that finds a solution, but then only the edges that appear in all input triangulations are forced to be true for all i , and again the edges that cross them are eliminated.

Furthermore, when computing a path, for every edge $uv \in T_\ell$, we eliminate flip variables that remove uv , i.e. $f(u, v, u_2, v_2, i)$ for all u_2, v_2, i . Similarly, for every edge $u_2v_2 \in T_0$, we eliminate flip variables that insert u_2v_2 , that is $f(u, v, u_2, v_2, i)$ for all u, v, i .

3.2 Crossing Lower Bound

Theorem 1 implies that, when building a path that starts at a triangulation T_0 , if an edge uv crosses $\chi(uv, T_0)$ segments of T_0 , then we only need to define the edge variable $e(u, v, i)$ for $i \geq \log_2(\chi(uv, T_0) + 1)$. The bound of $b(uv, T_0) \geq \lceil \log_2(\chi(uv, T_0) + 1) \rceil$ is tight when all the edges of T_0 that cross uv share a common endpoint (Figure 4). However, there are different ways in which the edges of T_0 may cross uv that may allow higher values of $b(uv, T_0)$.




Figure 4 A path of length 3 to insert an edge that had 6 crossings.

We consider estimations of $b(uv, T_0)$ based on the sequence of triangles that contain an upper or a lower edge (Figure 5). An edge uv that crosses $\chi(uv, T_0)$ segments of T_0 is translated into a *string* s of $\chi(uv, T_0) - 1$ symbols in the alphabet \uparrow, \downarrow , according to which side of uv contains the edge that is not crossed by uv in each triangle.




Figure 5 A path of length 4 to insert an edge that had 6 crossings. Each triangle is labeled and colored as containing an upper or a lower edge.

Let $b(s)$ be the number of steps to insert an edge $u'v'$ in a triangulation defined for a point set in convex position with the same string s and only containing the points u', v' , and the endpoints of the segments crossed by $u'v'$. We conjecture that flipping edges on the boundary of the polygon defined by the triangles that cross uv never reduces the number of flips needed to make uv crossing free. More precisely, we have the following conjecture.

Conjecture 4. *Let T be a triangulation of a point set P not containing a segment uv with endpoints in P . Let s be the string of uv in T . The number of parallel flips to insert uv in T is at least $b(s)$.*

197 We use a heuristic together with a depth first search exact solution to estimate $b(s)$.
198 Sometimes the heuristic makes errors, and the depth first search can only be used for relatively
199 small numbers of crossings, which implies that $b(s)$ is not always a lower bound and hence
200 the solution may no be optimal.

201 The heuristic is rather complicated and sometimes give incorrect bounds. To fix that, we
202 use an instance-independent preprocessing to compute tight values of $b(s)$ for strings with
203 up to 23? symbols and store only the values where the heuristic incorrectly computes $b(s)$ in
204 a file. This file is loaded by our solver and stored in a hash table.

205 3.3 Improving a Solution

206 Given a solution $P_1, \dots, P_{|\mathcal{T}|}$, we may improve it as follows. We choose a random path P_i
207 and use the SAT formulation to find a solution where the length of P_i is decremented and all
208 other lengths remain the same. We repeat this as often as necessary. Notice that the method
209 may converge to a locally optimal minimal list of lengths that is not globally optimal.

210 If the number of clauses is too large, there are two different approaches that we can take
211 (and they may be combined). We may force the new solution to be close to the original one
212 by only creating edge variables that cross few edges in the corresponding triangulation of the
213 previous solution.

214 We may also trim the solution to a certain radius r , by only rebuilding the portion of the
215 solution that is within r steps from the center. In this case, it is helpful to use MaxSAT first,
216 in order to reduce the number of unit-flips performed in the last steps as follows. Given a
217 path T_0, \dots, T_ℓ , we first use a MaxSAT solver to find the path of length ℓ that minimizes
218 the number of unit flips performed from $T_{\ell-1}$ to T_ℓ . The MaxSAT formulation is equal to
219 the SAT formulation with soft clauses $\neg f(u, v, u_2, v_2, \ell)$ for each last step flip variable. We
220 then find the path of length $\ell - 1$ that minimizes the number of unit flips performed from
221 $T_{\ell-2}$ to $T_{\ell-1}$. We continue this way for r steps.

222 3.4 Path Heuristic

223 The SAT formulation finds the shortest path connecting two triangulations fairly well, but it
224 may be too slow in some cases. We also designed a heuristic that produces reasonably short
225 paths quickly. We are given two triangulations T_0, T' and the goal is to find a short path
226 from T_0 to T' .

227 We use a greedy approach that iteratively obtains a triangulation T_{i+1} from a triangulation
228 T_i as follows. Let F denote the set of possible unit flips in T_i . More precisely, the elements
229 of F are pairs e, e' of edges such that a unit flip from T_i removes e and inserts e' . We then
230 build a graph $G(F)$ with vertex set F and edges between two unit flips that share a triangle.
231 Notice that the possible parallel flips correspond to independent sets in $G(F)$. We assign
232 a weight to the vertices as follows. The weight of a vertex e, e' is the number of edges in
233 T' crossed by e' minus the number of edges in T' crossed by e . Vertices of zero or negative
234 weight are eliminated.

235 We then proceed to greedily find an independent set I in $G(F)$. We iteratively add to I
236 the *unmarked* vertex of maximum weight, breaking ties by minimum degree. We then *mark*
237 all the vertices in the closed neighborhood of I . We repeat until all vertices are marked.

238 This iterative approach is repeated until we reach T' , which will happen in $O(n^2)$ flips
239 because of the following Theorem from [7].

240 ► **Theorem 5.** *Let T, T' be two triangulations of the same points set. If $T \neq T'$, then there
241 exists a unit flip that replaces an edge $e \in T$ by an edge e' such that e crosses more edges of
242 T' than e' does.*

243 We further improve the heuristic using the squeaky wheel paradigm [8]. Initially, we
244 assign weight 1 to every edge. We modify the definition of the weight of a flip as follows.
245 The weight of a flip e, e' is the sum of the weights of the edges in T' crossed by e' minus the
246 sum of the weights of the edges in T' crossed by e . When we reach triangulation $T_{i+1} = T'$
247 from a triangulation T_i we increment the weight of all edges in T' that are not in T_{i+1} . We
248 repeat the greedy algorithm with the new weights, keeping the best solution found. This
249 procedure is repeated multiple times.

250 3.5 Solution Heuristic

251 We use the following strategy to obtain reasonably good initial solutions that we may use
252 as a basis to improve with the aforementioned methods. We start by adding the Delaunay
253 triangulation as a candidate center. We then build other candidate centers by performing
254 flips starting from the Delaunay triangulation. Given an edge e and a triangulation T , let
255 $\chi(e, T)$ denote the number of segments of T crossed by e . We repeatedly perform unit flips
256 that remove an edge e and add an edge e' maximizing

$$258 \sum_{T \in \mathcal{T}} \chi(e)^p - \chi(e')^p,$$

259 as long as the value of the sum is positive. We add the triangulation we obtained to the set
260 of candidate center and repeat the process for a different value of p .

261 We calculate the paths from the candidate centers to each solution, building a solution
262 pool. For the next step it will be useful to have a small number of unit-flips in the last flip
263 of the path from the input triangulation to the center. To do that, we either use the greedy
264 heuristic or a MaxSAT formulation.

265 To improve the solution pool, we pick a solution from the pool and look at the triangulations
266 that are one flip away from the center in each path. For each such triangulation, we compute
267 the distance to the input triangulations and add the solution to the pool as before.

268 4 Results

269 In this section, we present the computational results that we obtained with our implementation
270 of the aforementioned algorithms and heuristics. In Section 4.1, we present the results on
271 computing short paths between two given triangulations. In Section 4.2, we present our
272 exact solver, with and without the happy edges conjecture. In Section 4.3, we present the
273 heuristics used to find solutions to the instances that we could not solve exactly.

274 The solvers were coded in C++ and compiled using gcc and use a single thread. During the
275 competition, they were executed on several Linux computers, either using GNU Parallel [12]
276 for local executions or Slurm [13] for cluster executions. It was very useful to have access
277 to machines with 128GB or more RAM to be able to solve large SAT formulations with
278 CaDiCal [5], which has been able to solve SAT instances with more than 50 million variables
279 and 500 million clauses. The time measurements on this paper have all been taken on a
280 AMD Ryzen 9 9900X CPU and ASUS TUF GAMING B650M motherboard with 128GB of
281 RAM running Fedora Core 43.

282 **4.1 Path Calculation**

283 Computing short paths between two given triangulations is a key component to obtain good
 284 solutions. Typically, these paths are computing with an input triangulation as one extreme,
 285 and a triangulation that makes a reasonably good center as the other extreme. In this section,
 286 we use the Delaunay triangulation as one extreme, because it is a well defined triangulation
 287 that makes a reasonably good (but not very good) center. Table 1 shows the length of the
 288 path and the running time of different heuristics and SAT solutions with and without the
 289 happy edges conjecture. The paths are calculated from the Delaunay triangulation to the
 290 first input triangulation of several instances. The SAT paths are obtained by first running
 291 the squeaky wheel heuristic in both directions, and then iteratively decreasing the path
 292 length. Notice that the running time of the heuristics is much smaller, while the result is
 293 rarely more than 1 unit away from the optimal distance, especially if we take the minimum
 294 of both directions. The squeaky wheel heuristic is run for at most 16 iterations, but stops
 295 earlier if the length increases from one iteration to the next.

| 296 | 297 | 298 | Greedy | | Greedy | | Squeaky | | Squeaky | | SAT | | SAT | |
|------|-----|-----|-----------|-----|-----------|------|-----------|-----|-----------|---------|-----------|--------|-------|-----|
| | | | forward | t | backward | t | forward | t | backward | t | happy | t | exact | t |
| 299 | 300 | 301 | 302 | 303 | 304 | 305 | 306 | 307 | 308 | 309 | 310 | 311 | 312 | 313 |
| 500 | 13 | 4.2 | 12 | 3.7 | 13 | 45 | 12 | 18 | 12 | 1512 | 12 | 10191 | . | . |
| 1000 | 12 | 7.8 | 12 | 6.8 | 12 | 25 | 12 | 38 | 11 | 6806 | 11 | 78264 | . | . |
| 1500 | 13 | 13 | 12 | 12 | 12 | 52 | 12 | 23 | 11 | 19302 | 11 | 201353 | . | . |
| 2000 | 15 | 21 | 13 | 18 | 13 | 65 | 13 | 104 | 13 | 16937 | 13 | 485177 | . | . |
| 3000 | 15 | 36 | 15 | 33 | 14 | 193 | 15 | 433 | 14 | 140729 | . | . | . | . |
| 4000 | 18 | 57 | 16 | 51 | 16 | 577 | 15 | 655 | 15 | 101402 | . | . | . | . |
| 5000 | 18 | 70 | 17 | 70 | 16 | 238 | 17 | 278 | 16 | 1037840 | . | . | . | . |
| 6000 | 17 | 93 | 16 | 83 | 17 | 1071 | 16 | 644 | 16 | 271081 | . | . | . | . |
| 7000 | 17 | 94 | 17 | 96 | 17 | 180 | 17 | 407 | 16 | 471389 | . | . | . | . |


308 **Table 1** Length and computation time (in milliseconds) from the Delaunay triangulation to the
 309 first input triangulation of the `rirs-n--20` instance for different values of the number of points n .
 310 The columns respectively correspond to the greedy heuristic forward and backward, the squeaky
 311 wheel heuristic forward and backward, the SAT solution with the happy edges conjecture, and the
 312 SAT solution without the happy edges conjecture, unless it takes too long. The best lengths found
 313 are shown in bold.

314 **4.2 Exact Solutions**

315 This section presents the exact solver that we used to solve most instances exactly. Figure 6
 316 shows the number of exact solutions found as a function of the running time, both without
 317 and with the happy edges conjecture (the value of the solutions found in both cases is the
 318 same, hence the conjecture stands). Notice that the impact of the happy edges conjecture is
 319 less significant than when computing only paths, as there are few common edges on all input
 320 triangulations. We remark that during the competition, we managed to find exact solutions
 321 to 189 of the 201 small instances without assuming the happy edges conjecture.


324 **4.3 Heuristic Solutions**

325 The typical process to solve instances that we did not solve exactly consists of several steps.
 326 First, we use the heuristic from Section 3.5 to obtain a reasonably good center. Notice that



322 ■ **Figure 6** Number of exact solutions found as a function of the running time over 3 hours of
323 execution with and without assuming the happy edges conjecture.


327 no SAT solver is used in this part. The evolution of the objective value for this step is shown
328 in Figure 7. The name of the `rirs` instances is composed of two values. The first is the
329 number of points and the second is the number of input triangulations.



330 ■ **Figure 7** Evolution of the best solution found by the heuristic solver without any SAT solver for
331 different instances.

332 Second, we build a new solution that keeps the same center but we recalculate the paths
333 using the sat formulation from 2.1, assuming the happy edges conjecture (Section 3.1) and
334 inexact lower bounds (Section 3.2). This will typically reduce the objective value by 1 to 4
335 percent. The calculation of 50 paths in each solution is shown in Figure 8. Notice that the
336 running time increases rapidly with the number of points and that finding a shorter path
337 (satisfiable SAT problem) is typically slower than when no shorter path exists (unsatisfiable
338 SAT problems).

342 Third, we improve the solution using the SAT formulation from Section 2.2, assuming
343 the happy edges conjecture (Section 3.1) and inexact lower bounds (Section 3.2). This part
344 requires adjusting a large number of parameters in order to obtain SAT problems that are
345 not too hard. The parameters include the distance to the previous center, the distance to
346 the previous path, and whether the solution will be trimmed. Optionally but recommended
347 when the problem is trimmed, we use MaxSAT to reduce the number of unit flips close to
348 the center. The improvement of some solutions over time is shown in Figure 9. Notice that



339 ■ **Figure 8** Path by path improvement of the heuristic solutions using a SAT solver, where each
 340 path is recalculated but the center is kept unchanged. Each circle represents the beginning of the
 341 computation of a path.

349 there is a significant preprocessing time to calculate edge variables with a limited number of
 350 crossings. This preprocessing is applied initially and after every improvement. Also notice
 351 that unsatisfiable SAT problems, which mean no improvement in the solution, are solved
 352 much faster than satisfiable ones.

358 5 Conclusion and Open Problems

359 We were surprised that we managed to solve so many instances exactly and how well an
 360 heuristic approximates the exact distance. We believe the following factors made the problem
 361 easy:

- 362 ■ The short path length that allows a somewhat small number of variables.
 363 ■ A SAT model where most clauses have size 2.
 364 ■ The ability to eliminate many variables through several arguments.
 365 ■ The fact that the happy edges conjecture is either true or at least holds in most practical
 366 cases.


367 The short path lengths are somewhat surprising, given that an $\Omega(n)$ lower bound is
 368 presented in [6], in contrast to the $\Theta(\log n)$ bound for combinatorial triangulations [9], where
 369 we can flip non-convex quadrilaterals. In the challenge instances, we observed path lengths
 370 that are roughly $O(\log n)$.

371 Still, there are instances with only 160 points and 20 input triangulations that we could
 372 not solve exactly and we still managed to improve solutions to instances with as few as 320
 373 points and 20 input triangulations months after the beginning of the challenge.

374 The challenge instances did not have many points in convex position. Surprisingly, the
 375 case where all points are in convex position is the hardest for our SAT formulation, as there
 376 are $\Theta(n^4)$ empty convex quadrilaterals. We wonder if a different model works better when
 377 all and also when most points are in convex position.

378 We believe that the same problem with unit flips is significantly harder because the long
 379 path lengths make the SAT formulation much more complex and removing a happy edge can
 380 reduce a path length from $\Theta(n^2)$ to $\Theta(n)$.

381 Many theoretical open problems remain, such as proving Conjectures 3 and 4. We also
 382 wonder if parallel flip distance problem is NP-hard in general and in convex position, as well



353 ■ **Figure 9** Improvement of the whole solution using a SAT solver and reducing the length of a
 354 random path by one unit at a time. We constrain the edges in the new solution to cross at most 3
 355 edges of the corresponding triangulation in the solution (except for the instance **rirs-500-20**, where
 356 the slower but more effective execution with the parameter set to at most 7 edges is also pictured)
 357 but do not use trimming. Each circle represents a new path length that we try to decrement.

383 as the unit flip distance problem in the convex case (the problem is NP-hard for general
 384 points sets [10]). We also do not know if the problem of calculating $b(s)$ (see Section 3.2)
 385 can be solved in polynomial time, possibly using dynamic programming.

386 ————— **References** —————

- 387 1 Oswin Aichholzer, Brad Ballinger, Therese Biedl, Mirela Damian, Erik D Demaine, Matias
 388 Korman, Anna Lubiw, Jayson Lynch, Josef Tkadlec, and Yushi Uno. Reconfiguration of
 389 non-crossing spanning trees, 2022. [arXiv:2206.03879](https://arxiv.org/abs/2206.03879).
- 390 2 Florent Avellaneda. A short description of the solver EvalMaxSAT. *MaxSAT Evaluation*,
 391 8:364, 2020.
- 392 3 Ruy Fabila-Monroy, Clemens Huemer, and Dieter Mitsche. The number of empty four-
 393 gons in random point sets. *Electronic Notes in Discrete Mathematics*, 46:161–168, 2014.
 394 doi:[10.1016/j.endm.2014.08.022](https://doi.org/10.1016/j.endm.2014.08.022).
- 395 4 Sándor P. Fekete, Phillip Keldenich, Dominik Krupke, and Stefan Schirra. Minimum coverage
 396 by convex polygons: The CG:SHOP challenge 2023, 2023. [arXiv:2303.07007](https://arxiv.org/abs/2303.07007).
- 397 5 Nils Froleyks, Marijn Heule, Markus Iser, Matti Järvisalo, and Martin Suda. CaDiCal, Kissat,
 398 Paracooba, Plingeling and Treengeling entering the SAT competition 2020. *Artif. Intell.*,
 399 301:103572, 2021. doi:[10.1016/j.artint.2021.103572](https://doi.org/10.1016/j.artint.2021.103572).

- 400 6 Jerôme Galtier, Ferran Hurtado, Marc Noy, Stéphane Pérennes, and Jorge Urrutia.
401 Simultaneous edge flipping in triangulations. *International Journal of Computational Geometry
402 & Applications*, 13(02):113–133, 2003.
- 403 7 Sabine Hanke, Thomas Ottmann, and Sven Schuierer. The edge-flipping distance of
404 triangulations. *Journal of Universal Computer Science*, 2(8):570–579, 1996.
- 405 8 David E. Joslin and David P. Clements. Squeaky wheel optimization. *Journal of Artificial
406 Intelligence Research*, 10:353–373, 1999.
- 407 9 Tanvir Kaykobad. An interlaced algorithm for transforming plane triangulations using
408 simultaneous flips. *Journal of Computational Geometry*, 16(1):517–550, 2025.
- 409 10 Alexander Pilz. Flip distance between triangulations of a planar point set is APX-hard.
410 *Computational Geometry*, 47(5):589–604, 2014. doi:10.1016/j.comgeo.2014.01.001.
- 411 11 Daniel D. Sleator and William P. Tarjan, Robert E. and Thurston. Rotation distance,
412 triangulations, and hyperbolic geometry. In *Proceedings of the eighteenth annual ACM
413 symposium on Theory of computing*, pages 122–135, 1986.
- 414 12 O. Tange. Gnu parallel - the command-line power tool. *:login: The USENIX Magazine*,
415 36(1):42–47, Feb 2011. URL: <http://www.gnu.org/s/parallel>.
- 416 13 Andy B Yoo, Morris A Jette, and Mark Grondona. Slurm: Simple linux utility for resource
417 management. In *Workshop on job scheduling strategies for parallel processing*, pages 44–60,
418 2003.